

Design Journal - Ethan Bennett

M1

1/2/2023 (~1 HR 30 M)

I would've liked to meet sooner, but with a teammate unable to access a computer all the way in Singapore it was quite difficult until now. The primary goal of this meeting was to figure out

- A) Register assignments
- B) Instruction types and format
- C) Processor architecture

So I guess I'll discuss those next.

First, we decided to do a Load/Store architecture, since we were the most confident in it. Had a lot of ideas for possible instruction formats right off the bat. We briefly considered an accumulator architecture, but some of the instruction types seemed really annoying, so we pivoted to Load/Store.

We agreed that we should probably have 16 registers, since that would give us more than enough. Something to consider is that having excess registers is quite expensive, but for now it's ok. Allocating them was easy enough, many similarities with RISC-V register allocation. Zero register, stack pointer, etc. We have significantly less save/temp/argument registers, though. One of the perks of 16 bit architecture is that having 32 registers is hard since that means any register-register operation is nearly impossible. That's like, 15 of 16 bits dedicated to just registers. Enjoy a one bit opcode. Luckily, we aren't psychopaths, and went with 16 registers. That leaves us room for a 3 bit opcode and a 1 bit funct.

On the subject of opcodes, our next priority was instruction types. Many of these share a lot of similarity with RISC-V, so it wasn't terribly difficult. The design document covers these quite nicely, so I won't bother repeating myself here. One thing we considered and had to scrap was merging the destination and source register in I-Type instructions (to allow for a larger immediate) but this meant we would need a whole new type for lw instructions. I didn't want to give lw the satisfaction of its own type. Thus, we went with a more traditional I-Type, as seen below.

I-Type (Immediate)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Imm[3:0]				rs1				rd				funct	opcode		

Allows for two immediate operations (adding a register value with an immediate, and loading a value to a register from memory).

Instruction	Opcode	Funct	Example	Meaning	Description
addi	001	0	addi x5, x6, 7	$x5 = x6 + 7$	$R[rd] = R[rs1] + SE(imm)$
lw	001	1	lw x5, 0(x6)	Load a value from memory address x6 with an offset of 0 bytes into x5.	$R[rd] = M[R[rs1]] + SE(imm)$

That's about all there is to say about this meeting.

1/9/2023 (4 HR)

Probably should have met sooner, but once again one of our teammates was still in Singapore until yesterday. Its fine, how hard could finishing the design document be-

4 HOURS LATER

OK, the design document is finished. Definitely will start work on M2 sooner. Geez. Had to do some rethinking on branch instructions, and realized we could just implement bne using beq, so we replaced bne with bgt. That's all really.

M2

Not a particularly interesting milestone to talk about in terms of big decisions. We cleaned up the design document a bit, and drew up the RDL. My main contribution was all of the component specifications - when doing that part, we had not yet considered multicycle, but that is something we should absolutely include since we are measuring performance based on execution time. We will need to add a few registers, but it shouldn't be excessively hard.

M3

The main concerns with this milestone were how to best transition to multicycle. Honestly? It's not that bad. We essentially just added a few registers in between the major components. The

RTL for the processor changed as a result, but nothing was unmanageable. I also drew up a prospective data path, but that aspect is not yet finalized.

M4

This one is a doozy. Here are the patch notes for my contributions

- Implementations of every component specified in the design document except for the program counter, register file, and mux
- I fixed the mux and register file, since they wouldn't compile
- Except for the program counter and 2:1 mux, I have written every functional test bench
- Added the datapath (no one else has submitted an alternative, despite us agreeing to do so, I guess my word is gospel)
- Drew the control state diagram (flawed, deeply flawed, will be corrected in M5)
- Began integration testing a la design document plan

Note we are still missing the **branch** control signal, because I forgot it.

M5

Most important correction from last time is that the control state diagram is not completely insane anymore. Still no branch signal - I am but one man - but it is much better. Another new thing is a much better integration plan than some words that don't mean a whole lot, refer to the design document if you are curious. I tried my hand at integrating more parts, but to be frank, that is far too much work for me alone. We desperately need to allocate tasks better. We essentially haven't finished this milestone, after all.

M6

Integration, integration, integration, that's the name of the game. Frankly, I did pretty well. The only problem is the thing doesn't work. We almost certainly will not finish it, but we narrowed it down to the register file and control unit - the former wasn't able to be written to (how did no one catch that?) and the latter had timing issues that causes our instructions to not work. At this point, we essentially made the call to focus on making our design document as professional as possible; I mostly supplied testing results and integration plan stuff. Mac is the real MVP here; really came through on this design doc, as well as the assembler.

Retrospectively, I really wish we went with less registers. I also wish we spent more time in general (though I know we easily spent 100ish man hours on this) but getting started earlier so we could ask more questions would've been nothing but a benefit.